

Annotation of Emotion Carriers in Personal Narratives

Aniruddha Tammewar¹, Alessandra Cervone¹, Eva-Maria Messner², Giuseppe Riccardi¹

¹Signals and Interactive Systems Lab, University of Trento

²Clinical Psychology and Psychotherapy, University of Ulm

{aniruddha.tammewar, alessandra.cervone, giuseppe.riccardi}@unitn.it, eva-maria.messner@uni-ulm.de

Abstract

We are interested in the problem of understanding personal narratives (PN) - spoken or written - recollections of facts, events, and thoughts. In PN, emotion carriers are the speech or text segments that best explain the emotional state of the user. Such segments may include entities, verb or noun phrases. Advanced automatic understanding of PNs requires not only the prediction of the user emotional state but also to identify which events (e.g. *the loss of relative* or *the visit of grandpa*) or people (e.g. *the old group of high school mates*) carry the emotion manifested during the personal recollection. This work proposes and evaluates an annotation model for identifying emotion carriers in spoken personal narratives. Compared to other text genres such as news and microblogs, spoken PNs are particularly challenging because a narrative is usually unstructured, involving multiple sub-events and characters as well as thoughts and associated emotions perceived by the narrator. In this work, we experiment with annotating emotion carriers from speech transcriptions in the Ulm State-of-Mind in Speech (USoMS) corpus, a dataset of German PNs. We believe this resource could be used for experiments in the automatic extraction of emotion carriers from PN, a task that could provide further advancements in narrative understanding.

Keywords: Personal Narratives, Emotion Analysis, Sentiment Analysis

1. Introduction

Sentiment and emotion recognition is a well-grounded research field in the NLP community. People express their emotions directly or indirectly through their speech, writings, facial expressions or gestures. There has been work on analyzing sentiment using various modalities such as speech, text, visuals, biosignals, and combinations of these (Poria et al., 2016; Kudiri et al., 2016; Soleymani et al., 2015). In text-based communication, there are various mediums through which people express their emotions, such as social media, personal diaries, news articles, blogs. In conventional emotion and sentiment analysis systems, emotion, in general, represents a strong feeling deriving from the circumstances whereas the sentiment is the opinion expressed based on the feelings. The sentiment is thus the effect of the emotion (Sailunaz et al., 2018). While sentiment is measured by a numeric score ranging *negative* to *neutral* to *positive*, basic emotions are usually classified into classes such as *anger*, *disgust*, *fear*, *joy*, *sadness*, *surprise* using established emotion annotation schemes (Ekman, 1992; Shaver et al., 1987; Oatley and Johnson-Laird, 1987).

While most research on emotion analysis focuses on emotion classification, including predicting emotions of the writer as well as those of the reader of a text (Chang et al., 2015), a few studies have focused on identifying what might have triggered that emotion (Sailunaz et al., 2018). In particular, the task has been framed as *Emotion Cause Extraction* (Lee et al., 2010a) from news and microblogs, where the cause of an emotion is usually a single clause (Chen et al., 2010) connected by a discourse relation to another clause that explicitly expresses a given emotion (Cheng et al., 2017), as in this example from Gui et al. (2016a): “< cause > Talking about his honours, < /cause > Mr. Zhu is so < emotion > proud < /emotion >.”

However, if we move to other text genres beyond news and microblogs, it might be more complicated to identify the

causes of given emotions from the text. In this work, we focus on spoken personal narratives. A Personal Narrative (PN) can be defined as a recollection of an event or connected sequence of events the narrator has been part of, as an active or passive participant. Examples of PNs include personal diaries, short notes, travelogues in digital form (speech and/or text). Compared to other text genres such as news and microblogs, spoken PNs are particularly challenging because the structure of the narrative could be complex, involving multiple sub-events and characters as well as thoughts and associated emotions perceived by the narrator.

Tammewar et al. (2019), for example, observed that in different machine learning models (Support Vector Machine, Attention-based neural sequence tagger) trained to predict valence (emotional value associated with a stimulus) from spoken PNs, concepts beyond sentiment words (*sad*, *happy*) were found to be useful. These concepts included terms such as characters (e.g. *grandfather*, *a friend*), locations (e.g. *swimming pool*) and events (e.g. *high school exam*). As the task of the models was to predict the emotional state of the narrator, the authors concluded that these concepts played the role of explaining and *carrying* the emotional state of the person. We call such concepts ‘*emotion carriers*’.

Inspired by such evidence, in this work, we investigate the possibility of annotating emotion carriers in spoken PNs. This type of annotation could then be used to experiment with building automatic Emotion Carriers Extraction systems. Such models, together with emotion prediction models, could provide a deeper understanding of the emotional state of the narrator. The task of emotion carriers extraction can be classified as a task of Automatic Narrative Understanding (ANU), which encompasses tasks that extract various information from narratives (Fu et al., 2019). Emotion carriers extraction, for example, could be a useful task for conversational mental healthcare applications. Applications aimed at the mental well-being of the users, often col-

German (original)	English (translated)
<p>Okay. Ähm also eine Situation, in der ich mich kompetent gefühlt hab, war, als ich meinen ähm Praktikumsplatz bekommen hab und ähm ich in dem Praktikum dann auch ähm Sachen selbstständig machen durfte und auch die Rückmeldung bekommen hab von den Personen dort, dass das, was ich da so mach, dass das gut ist und dass ähm sie mit mir sehr, mit mit mir sehr zufrieden sind ähm. Und die Gefühle dabei waren natürlich irgendwie Glück, weil man ist davor unsicher, ob man das, was man da macht, ob das so gut ist und ob man das so schafft. Ähm und das hat eben sehr gut funktioniert. Also ich hab mich sehr zufrieden gefühlt, mit mir ähm im Reinen, mit mir glücklich, auch irgendwie so ein bisschen Bestätigung darin bekommen, dass das, was ich mach, gut ist oder das, was ich auch jetzt als Studium gemacht hab, irgendwie passt. Ähm ähm so bisschen so positive Aufregung, also man fühlt sich sehr sehr wach, erregt irgendwie, aber in einer positiven Art und Weise. ...</p>	<p>OK. Um, so a situation in which I felt competent, was when I got my um internship position and er in the internship then I was also allowed to do things independently and also got the feedback from the people there, that what I am doing there, that it is good, and that they are very pleased with me, with me, um. And of course the feelings were kind of lucky, because you are not sure if you, what you are doing, if that is so good and if you can do it that way. Um, and that worked very well. So I felt very satisfied, with me uh, I'm happy, with me, somehow getting a bit of confirmation that what I'm doing is good or what I'm doing now as a study have, somehow fits. Uhm umh so a bit so positive excitement, so you feel very very awake, excited somehow, but in a positive way. ...</p>

Table 1: **PN Annotated Example** (Colored Figure). A part of a narrative showing text spans annotated by the four annotators. The intensity of the red color in the background represents the number of annotators who annotated the text-span (varying from lightest for 1 annotator to the darkest for four annotators). It can be seen that some spans are annotated by single annotator while some by multiple. In the text-span “*positive Aufregung*”, two annotators selected the entire span while another one selected only the second word “*Aufregung*”. The annotations contain sentiment words as well as content words.

lect personal narratives from users in the form of personal diaries. In this setup, the applications may benefit from an analysis that provides not only the user’s emotion trends but also the emotion carriers. The conversational agent of the application could use this information to start a conversation with the user and elicit more information about them, which in turn can be provided to a therapist for personalized interventions.

Nevertheless, annotating emotion carriers in PNs is a challenging task, since the same term may or may not carry emotions in different contexts and thus the search space for identifying emotion carriers becomes huge.

In this work, we present our experiments with manual annotation of emotion carriers in spoken PNs in the German language from the Ulm State-of-Mind in Speech (USoMS) corpus (Schuller et al., 2018). After a review of relevant existing tasks in Section 2., we describe the USoMs corpus in Section 3.. The challenge of the task of Emotion Carriers extraction, formalized as the extraction of text segments perceived to be crucial for predicting the user’s emotional state, is described in Section 4., while the details about the manual annotation experiments are presented in Section 5.. In Section 6. we analyze and evaluate the annotated data using inter-annotator agreement metrics and qualitative analysis to get insights about the complexity of the task. Finally, in Section 7., we draw on the conclusions of our work.

2. Related Work

While the Emotion Carriers in PNs is a new field of research, the most relevant studies have worked on the task of *Emotion Cause Extraction (ECE)*. The task focuses on

finding the cause of emotion from the given text. According to (Talmy, 2000), the cause should be an event itself. The cause-event refers to the immediate cause of the emotion, which can be the actual trigger event or the perception of the trigger event (Lee et al., 2010a). The cause events are further categorized into two types: verbal events and nominal events (Lee et al., 2010b).

There are a few emotion cause corpora on formal texts such as news reports, frames from FrameNet (Tokuhisa et al., 2008; Ghazi et al., 2015; Lee et al., 2010a; Gui et al., 2016a) and informal texts such as microblogs (Gui et al., 2014; Gao et al., 2015b; Cheng et al., 2017). Some of these are annotated manually, while some are created automatically.

While Lee et al. (2010a) defined the task of ECE as the extraction of word-level emotion causes, Chen et al. (2010) suggested that a clause would be a more appropriate choice of unit for the extraction of an emotion cause. There have been works trying to solve the problem using different methods: Rule-based (Neviarouskaya and Aono, 2013; Li and Xu, 2014; Gao et al., 2015b; Gao et al., 2015a; Yada et al., 2017); Machine Learning based (Ghazi et al., 2015; Song and Meng, 2015; Gui et al., 2016a; Gui et al., 2016b; Xu et al., 2017) and Deep Learning based (Gui et al., 2017; Li et al., 2018; Yu et al., 2019; Xu et al., 2019).

While most of the previous works are focused on either the news domain or microblogs, our work focuses on personal narratives. Personal narratives are more complex than the other domains as they are typically longer and contain multiple sub-events. Moreover, each sub-event has attributes (such as characters, entities involved in the sub-event), the

narrator’s reactions and emotions expressed in the narrative. In such a complex sequence of sub-events, it is difficult to associate the emotion clause with the corresponding event. It was one of the main shortcomings in the work by Gui et al. (2016a), they call it the problem of *cascading events*. Even if we succeed in correctly extracting the cause clauses for each sub-event, it is not our final goal. Our goal is to extract the emotion carriers used for conveying the emotions manifested in the narrator after recounting the entire narrative/event. The emotions produced by sub-events may or may not represent the emotions at the level of the entire narrative and thus the carriers as well. For instance, a narrative could begin with a happy event that is told to build context but could end up in a sad situation. Note that although the microblogs could also be considered as a personal narrative, we are dealing with longer personal narratives.

The above approaches make one important assumption that an emotion keyword is always present in the text, for which they have to find the cause (Gui et al., 2016a). As described by (Cheng et al., 2017), they consider the task of ECE as a discourse relation between the cause-clause and the emotion-clause. However, personal narratives may or may not contain emotion keywords. They might contain multiple keywords as well. Particularly in this task, the dataset is provided with the sentiment (positive or negative) of the text, and our goal is to find the emotion carriers for that sentiment.

3. USoMs Corpus

Ulm State-of-Mind in Speech (USoMs) is a database of spoken PNs in German, along with the self-assessed valence and arousal scores. A part of the dataset was used and released in the Self-Assessed Affect Sub-challenge, a part of the Interspeech 2018 Computational Paralinguistics Challenge (ComParE) (Schuller et al., 2018). The task was to predict the narrator’s valence score provided a short speech fragment (8 seconds) of the narrative.

The data consists of 100 speakers (students) (85 f, 15 m, age 18-36 years, mean 22.3 years, std. dev. 3.6 years). In the challenge, the data was divided into three sets *training, development, and test*. In this study, we annotate the development and the test sets (66 participants’ data). The students told two negative and two positive PNs, each with a duration of about 5 minutes. Before and after recording each narrative, they self-assessed valence (spanning from negative to positive) and arousal (spanning from sleepy to excited) using the affect grid (Russell, 2003) on a 10-point Likert scale. The narratives were transcribed manually. The number of tokens in vary from 292 to 1536 (mean: 820; std: 208). We use these transcripts in our work to enrich the annotations with the emotion carriers.¹

Following prompts were used to elicit the narratives 1) Negative narrative: *"Please remember a time in your life when you were facing a seemingly unsolvable problem and report as detailed as possible over the next five minutes"*. 2) Positive narrative: *"Please report of a time in your life were*

you found a solution, where you felt powerful, happy, and content. Describe that story in-depth over the next five minutes".

4. The challenge of Emotion Carriers Extraction in spoken Personal Narratives

While describing an emotional event, to convey the emotions narrators not only use explicit emotion words such as *happy, sad, excited* but also other emotion carriers such as entities, persons, objects, places, sub-events related to the event. (Tammewar et al., 2019), in their work, found that these carriers play an important role in predicting the current mental state of the narrator. There could be mentions of many such carriers in the narrative, but this does not imply that all of them reflect/carry the narrator’s emotions. We call the text spans that could capture the annotator’s emotion, *"emotion carriers"*.

Apart from the carriers themselves, other factors also influence their importance, such as a context (from the narrative), the position of the span in the narrative, or the frequency of the mentions. Consider the term *grandfather*. The emotional value/valence associated with the term changes according to the provided context, it can carry emotion accordingly. Consider the phrases *"my grandfather died"* and *"my grandfather came to visit us after a long time"*. In the first case, it might carry a negative emotion while, in the second case, a positive emotion. Often, narratives are longer, and people mention some terms to build the context for the main event. Consider the example *"That day, my grandfather had come to visit us after a long time... we all went to the beach, where a saw a dreadful accident. A boy was swept out to sea while walking on the outer banks with his mom.... since then, I'm afraid of the water bodies."* Here, the term *grandfather* is used just to ground the context, while the emotion-relevant linguistic carriers come later in the narrative, such as *"swept out"*, *"water bodies"*. This uncertainty makes the task of identification of emotion carrying terms, complex and subjective.

5. Annotation of Emotion Carriers

To investigate the possibility of identifying emotion carriers from personal narratives, we annotate the personal narratives from the USoMs Corpus explained in Section 3., with the text spans that carry the emotions of the narrator manifested during the recollection.

In our annotation scheme, even though it is more relevant to us, we do not provide the annotators with a pre-selected noun or verb phrases to choose from. We give them the freedom to select text segments they feel are most important for our task. We believe that the pre-selection of spans could build bias in annotators towards specific fragments while there could be other text-fragments which are more important emotion carriers. Also, being spoken narratives, the automated tools to extract the noun and verb phrases may produce errors, thus affecting the annotation quality.

In this section, we provide details of the annotation experiments, including the annotators, the annotation scheme, and a brief overview of the tool used for the annotation.

¹We plan to organize a challenge and release the annotated data as a part of it

5.1. Annotators

Each narrative is annotated by four annotators. All the annotators are native German speakers and hold a Bachelor's degree in Psychology. They have been specifically trained to perform the task. We refer to the four annotators 'ann1', 'ann2', 'ann3' and 'ann4'

5.2. Annotation Guidelines

The annotation task involves the selection of the emotion carrying text spans as perceived by the annotator. We provide annotators with the guidelines to follow while performing the task.

We ask them to select sequences of adjacent words (one or more) in the text that explain why the narrative is positive or negative for the narrator.

We are particularly interested in words that play an important role in the story, such as:

- People (e.g. 'mother', 'uncle John', 'my best friend'); Locations ('university', 'our old school'); Objects (e.g. 'guitar', 'my first computer'); Events ('exam', 'swimming class', 'prom night')
- A clause that can include a verb and nouns (e.g. 'Mary broke my heart', 'I lost my guitar', 'I failed the admission exam')

They have to select a minimum of three such text spans.

We also provide them with the best practices to be followed:

- We ask them to annotate the contentful words ('university', 'mother') preferably over pronouns ('she', 'her', 'it')
- If the same term is present multiple times, they are asked to annotate the first instance of the same concept and to avoid repetition.
- To make sure if something needs to be added or removed from the list of selected fragments, the annotators are asked to make sure:
 - If a person who has not read the narrative can understand why the event was positive or negative just by looking at the list of spans they have selected. If not, they have to check if something is missing.
 - They are asked to ensure that there are no repetitions in the list and that there are no spans, which are not central to the narrative.
- As the annotators already know if the narrative is positive or negative in general, we ask them to annotate the feelings (emotion words) only if they are more informative (e.g. 'feeling of freedom') than simple positive/negative (e.g. 'I was happy').

5.3. Annotation Tool

We provide the annotators with a web-based tool to perform the annotations. The tool is mainly divided into two parts. In one part, we show them a personal narrative and the corresponding sentiment. The annotator can hover over

the tokens and select text spans by clicking and dragging over the consecutive tokens. On the right-hand side, they can see the already selected spans and their rankings. They can change the ranking by simple drag and drop.

Table 1 shows an example of annotations for a part of a narrative. We observe that sometimes, annotators annotate text-segments representing a similar concept but are at different positions in the text. In the example, the terms *Praktikumsplatz* and *Praktikum* represent the same concept of *internship* but two of the annotators followed the guidelines to select the first occurrence while the other annotator selected the second occurrence of the same concept.

6. Analysis

In this section, we perform an analysis of the annotations performed on the USoMs corpus from different perspectives. We take a look at some statistics of the annotations then we evaluate the annotations by calculating inter-annotator agreements with different strategies and finally we discuss some important observations we made.

6.1. Statistics

In this study, we analyze 239 narratives from 66 participants (the development and test sets from the ComParE challenge) that have been annotated by four annotators each. Note that for 66 participants the total number of narratives should be 264, but in the ComParE challenge, 25 files were removed because of issues like noise.

We observe that the number of annotations (text-spans) annotated by the annotators per narrative vary from 3 to 14 with an average of 4.6, also that all annotators follow the same pattern from this aspect. We also calculated the number of tokens present in the annotations. The numbers show that three of the annotators (*ann1*, *ann2*, *ann3*), on average select a span of 1.5 tokens, while the fourth annotator (*ann4*) selects three tokens (avg.) per annotation. Note that, for all the analysis, we use the spaCy toolkit² for tokenization. We observe that many annotations contain punctuation marks, which are considered as separate tokens by spaCy. Thus, we perform the same calculations while ignoring the punctuation tokens. We find that the average number of tokens drops down to 1.1 for the first three annotators, while it drops down to 2.3 for the *ann4*.

Another observation we make is that the annotations are scattered across all positions in the narratives. The mean position of the annotations is usually near the middle of the narratives. Apart from these aspects, we also analyzed the distributions of POS tags, and as expected, found that the top four categories include adjectives, nouns, verbs, and adverbs.

6.2. Evaluation of Annotations

Commonly used metrics for evaluating the agreement between annotators include variations of κ coefficient (a chance-corrected percent agreement measure) such as Cohen's (Cohen, 1960) for two annotators, Fleiss' (Fleiss, 1971) for multiple annotators. Unfortunately, calculations for κ such as observed and chance agreements involve the

²<https://spacy.io/>

	ann1	ann2	ann3	ann4
ann1	1	0.344	0.417	0.125
ann2		1	0.389	0.106
ann3			1	0.137
ann4				1

(a) Exact match, position agnostic, token level
(mean F_1 : 0.252)

	ann1	ann2	ann3	ann4
ann1	1	0.338	0.42	0.277
ann2		1	0.381	0.196
ann3			1	0.308
ann4				1

(b) Partial match with position, token level
(mean F_1 : 0.320)

	ann1	ann2	ann3	ann4
ann1	1	0.397	0.483	0.402
ann2		1	0.439	0.264
ann3			1	0.404
ann4				1

(c) Partial match, position agnostic, token level
(mean F_1 : 0.399)

	ann1	ann2	ann3	ann4
ann1	1	0.400	0.490	0.410
ann2		1	0.440	0.267
ann3			1	0.413
ann4				1

(d) Partial match; position agnostic; lemma level (mean F_1 : 0.403)

Table 2: Pairwise Inter-Annotator Agreement scores (F_1 measure) with respect to the different matching strategies. We vary the matching criteria for annotations based on three aspects 1) checking if the annotations are exactly same (exact match) or calculating the overlap between them (partial match) 2) positions of the annotations in the text is considered while matching (with position) or not (position agnostic) and 3) to calculate the overlap, the tokens in the annotations are matched (token level) or the lemmas of the tokens are matched (lemma level).

knowledge of *true negatives*, which is not well defined for a text span selection task. (eg. in this study, it could mean the number of possible text spans that are not annotated). This makes κ impractical as a measure of agreement for text spans annotation.

An alternative agreement measure that does not require the knowledge of *true negatives* for its calculations is Positive (Specific) Agreement (Fleiss, 1975), which is similar to the widely used F-measure (Hripcsak and Rothschild, 2005). It has previously been shown to be useful in the evaluation of crowdsourced annotations tasks, similar to our’s (Stepanov et al., 2018; Chowdhury et al., 2014).

Another problem we face in the task of text spans selection is the annotation of overlapping text fragments. Given the freedom on the lengths and positions of the text spans, two annotators might annotate different but overlapping text spans. The overlapping part could be an important part, thus the annotations should not be discarded completely. For instance, in Table 1, ‘*positive Aufregung*’ and ‘*Aufregung*’, both the spans contain the fragment ‘*Aufregung*’, which is important to be considered. Thus, we report results on exact matches as well as partial matches, following the work by (Johansson and Moschitti, 2010). For the partial match, they calculate ‘*soft*’ F-measure by calculating the coverage of the hypothesis spans.

As the personal narratives are longer, often some terms are repetitive. In our task, the position of an annotation is not quite important compared to the content. We further try to loosen the criteria for matching by not considering the position of the text fragments. For instance, let us say a narrative contains mentions of ‘*trip*’ at multiple places, like ‘*we went for a trip to India*’ and ‘*the trip was great*’. If two annotators intend to annotate the word ‘*trip*’, they have multiple locations to choose from. While from the perspective of discourse, it would be interesting to see which position seems more appropriate, for our purpose of extrac-

tion of emotion carriers it is less important. Following the same intuition, we also try to match tokens having the same lemma.

Table 2 shows the evaluation results based on the various strategies of matching described above. The F-measure is calculated for all pairs of annotators. For each strategy, we also report the mean of pairwise scores. In the four tables from Table(a) to Table(d) we loosen the matching criteria, thus increasing the scores. We show the results starting from the most strict criteria of exact matching in the table (a), then in the table (b), we show results for partial matching, but the positions of the annotations are taken consideration. The improvements are most significant in the case of *ann4*, as we saw earlier in Section 6.1. that *ann4* usually annotates longer fragments than others. This shows that the *ann4* annotates longer spans, but still contains the important part that other annotators annotate. Later in table (C), we remove the constraint of position, which results in improved scores, showing that even if the annotations by different annotators are different they often contain similar terms/carriers. This also shows that the annotators often ignore the instruction from the guidelines of selecting the first occurrence of the same term (Section 5.2.). In the table (d), we further try to match more things by considering lemmas instead of tokens, which results in an increment.

6.3. Discussion

It is difficult to judge the quality of the annotations by looking at the inter-annotator agreement scores, which are not self-explanatory. If we compare our task with other previous tasks that used a similar metric, we can better understand the complexity of the task and judge the quality of the annotations. Chowdhury et al. (2014) worked on the task of semantic annotations of the utterances from conversations. One of the sub-task annotators had to perform was selecting a text span describing a hardware concept. For a particular




Figure 1: Count histogram of emotion carriers shared and agreed upon by the annotators. Notice the long tail having only one occurrence counts.

concept like *printer*, the annotators could select the spans ‘with the printer’, ‘the printer’ or just ‘printer’, all of which are correct. The problem they faced for the selection of span is similar to ours, but the complexity and subjectivity are low, as they work on shorter texts and the annotator has more idea about the concept to be selected. They use the same metric as ours to evaluate the inter-annotator agreement for the span selection. They achieve $f1$ scores of 0.39 and 0.46 (for two different subsets of data) for the exact match, whereas 0.63 and 0.7 for the partial match (mean of pairwise agreements between three annotators). While our scores for exact and partial matches are 0.25 and 0.4, which we believe are reasonable given the subjectivity of the task and more number of annotators.

We calculate the sentiment polarity of each annotation using the *textblob-de* library³, which simply makes use of the polarity scores of the words from *senti-wordnet* for German (with simple heuristics), similar to the English *senti-wordnet* (Esuli and Sebastiani, 2006). We find that the trends of using sentiment carrying phrases vary across the annotators. The fraction of annotations carrying sentiment varies from 24% to 56% (*ann1*: 39%; *ann2*: 24%; *ann3*: 36%; *ann4*: 56%) for the four annotators. For further analysis, we plan to categorize the annotations into categories inspired by the ones used in the *Psychological Processes* categories of the LIWC dictionary (Pennebaker et al., 2015).

In Figure 1, we show an interesting observation from the annotations. We plot the histogram of overlaps from partial-matches that we get while evaluating the inter-annotator agreements for all annotator-pairs, with respect to their counts of occurrences. (Note that, due to space limitations, we only show the representative annotations and

not all the partial-matches). We can see that the overlaps contain both, the emotion words such as *proud*, *hopeless*, *disbelief* as well as content words like *scholarship*, *education*, *dance*. Also, we notice that there is a long tail of overlaps having only a single occurrence. This shows that a few terms are annotated frequently and agreed upon by annotators, while many terms are unique to specific narratives.

7. Conclusion and Future Directions

We proposed a new annotation scheme for the task of extracting emotion carriers from personal narratives (PN), which provides a deeper emotion analysis compared to the conventional emotional prediction task. We performed manual annotations of personal narratives to extract the emotion carriers that best explain the emotional state of the narrator. The annotation was done by four annotators. Narratives being longer and having a complex structure, we find the task to be subjective, which is reflected in the inter-annotator agreement scores and other analyses. Nevertheless, we find surprisingly high overlaps over the annotations, consisting of content words.

As it is difficult to interpret the agreement-scores in itself, to judge the quality of the annotations, we plan to study the improvements in the end-applications making use of the annotations, such as the task of valence prediction. If using the annotations provides better results than the existing systems, we can conclude that the annotations are indeed useful.

We believe that automated extraction of the emotion carriers as a task of Automatic Narrative Understanding (ANU) could benefit various applications. For instance, a conversational agent could use this information (from a PN shared by user) to start a meaningful conversation with the user

³<https://textblob-de.readthedocs.io/en/latest>

about the extracted emotion carriers rather than just showing *sympathy* or *happiness* based on the emotion classification of the PN. We plan to annotate the training data and build a module for the automatic extraction of emotion carriers.

Another interesting aspect to study is the correlation between speech and emotion carriers. The USoMs corpus contains speech and transcriptions of the PNs, allowing us to explore the correlation between the different modalities

8. Acknowledgements

The research leading to these results has received funding from the European Union – H2020 Programme under grant agreement 826266: COADAPT.

9. Bibliographical References

- Chang, Y.-C., Chen, C.-C., Hsieh, Y.-L., Chen, C. C., and Hsu, W.-L. (2015). Linguistic template extraction for recognizing reader-emotion and emotional resonance writing assistance. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 775–780.
- Chen, Y., Lee, S. Y. M., Li, S., and Huang, C.-R. (2010). Emotion cause detection with linguistic constructions. In *Proceedings of the 23rd International Conference on Computational Linguistics*, pages 179–187. Association for Computational Linguistics.
- Cheng, X., Chen, Y., Cheng, B., Li, S., and Zhou, G. (2017). An emotion cause corpus for chinese microblogs with multiple-user structures. *ACM Transactions on Asian and Low-Resource Language Information Processing (TALLIP)*, 17(1):6.
- Chowdhury, S. A., Ghosh, A., Stepanov, E. A., Bayer, A. O., Riccardi, G., and Klasinas, I. (2014). Cross-language transfer of semantic annotation via targeted crowdsourcing. In *Fifteenth Annual Conference of the International Speech Communication Association*.
- Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and psychological measurement*, 20(1):37–46.
- Ekman, P. (1992). An argument for basic emotions. *Cognition & emotion*, 6(3-4):169–200.
- Esuli, A. and Sebastiani, F. (2006). Sentiwordnet: A publicly available lexical resource for opinion mining. In *LREC*, volume 6, pages 417–422. Citeseer.
- Fleiss, J. L. (1971). Measuring nominal scale agreement among many raters. *Psychological bulletin*, 76(5):378.
- Fleiss, J. L. (1975). Measuring agreement between two judges on the presence or absence of a trait. *Biometrics*, pages 651–659.
- Fu, L., Chang, J. P., and Danescu-Niculescu-Mizil, C. (2019). Asking the right question: Inferring advice-seeking intentions from personal narratives. *arXiv preprint arXiv:1904.01587*.
- Gao, K., Xu, H., and Wang, J. (2015a). Emotion cause detection for chinese micro-blogs based on ecocc model. In *Pacific-Asia Conference on Knowledge Discovery and Data Mining*, pages 3–14. Springer.
- Gao, K., Xu, H., and Wang, J. (2015b). A rule-based approach to emotion cause detection for chinese micro-blogs. *Expert Systems with Applications*, 42(9):4517–4528.
- Ghazi, D., Inkpen, D., and Szpakowicz, S. (2015). Detecting emotion stimuli in emotion-bearing sentences. In *International Conference on Intelligent Text Processing and Computational Linguistics*, pages 152–165. Springer.
- Gui, L., Yuan, L., Xu, R., Liu, B., Lu, Q., and Zhou, Y. (2014). Emotion cause detection with linguistic construction in chinese weibo text. In *Natural Language Processing and Chinese Computing*, pages 457–464. Springer.
- Gui, L., Wu, D., Xu, R., Lu, Q., and Zhou, Y. (2016a). Event-driven emotion cause extraction with corpus construction. In *EMNLP*, pages 1639–1649. World Scientific.
- Gui, L., Xu, R., Lu, Q., Wu, D., and Zhou, Y. (2016b). Emotion cause extraction, a challenging task with corpus construction. In *Chinese National Conference on Social Media Processing*, pages 98–109. Springer.
- Gui, L., Hu, J., He, Y., Xu, R., Qin, L., and Du, J. (2017). A question answering approach for emotion cause extraction. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 1593–1602.
- Hripcsak, G. and Rothschild, A. S. (2005). Agreement, the f-measure, and reliability in information retrieval. *Journal of the American Medical Informatics Association*, 12(3):296–298.
- Johansson, R. and Moschitti, A. (2010). Syntactic and semantic structure for opinion expression detection. In *Proceedings of the Fourteenth Conference on Computational Natural Language Learning*, pages 67–76. Uppsala, Sweden, July. Association for Computational Linguistics.
- Kudiri, K. M., Said, A. M., and Nayan, M. Y. (2016). Human emotion detection through speech and facial expressions. In *2016 3rd International Conference on Computer and Information Sciences (ICCOINS)*, pages 351–356. IEEE.
- Lee, S. Y. M., Chen, Y., and Huang, C.-R. (2010a). A text-driven rule-based system for emotion cause detection. In *Proceedings of the NAACL HLT 2010 Workshop on Computational Approaches to Analysis and Generation of Emotion in Text*, pages 45–53. Association for Computational Linguistics.
- Lee, S. Y. M., Chen, Y., Li, S., and Huang, C.-R. (2010b). Emotion cause events: Corpus construction and analysis. In *LREC*.
- Li, W. and Xu, H. (2014). Text-based emotion classification using emotion cause extraction. *Expert Systems with Applications*, 41(4):1742–1749.
- Li, X., Song, K., Feng, S., Wang, D., and Zhang, Y. (2018). A co-attention neural network model for emotion cause analysis with emotional context awareness. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 4752–4757.

- Neviarouskaya, A. and Aono, M. (2013). Extracting causes of emotions from text. In *Proceedings of the Sixth International Joint Conference on Natural Language Processing*, pages 932–936.
- Oatley, K. and Johnson-Laird, P. N. (1987). Towards a cognitive theory of emotions. *Cognition and emotion*, 1(1):29–50.
- Pennebaker, J. W., Boyd, R. L., Jordan, K., and Blackburn, K. (2015). The development and psychometric properties of liwc2015. Technical report.
- Poria, S., Chaturvedi, I., Cambria, E., and Hussain, A. (2016). Convolutional mkl based multimodal emotion recognition and sentiment analysis. In *2016 IEEE 16th international conference on data mining (ICDM)*, pages 439–448. IEEE.
- Russell, J. A. (2003). Core affect and the psychological construction of emotion. *Psychological review*, 110(1):145.
- Sailunaz, K., Dhaliwal, M., Rokne, J., and Alhajj, R. (2018). Emotion detection from text and speech: a survey. *Social Network Analysis and Mining*, 8(1):28.
- Schuller, B. W., Steidl, S., Batliner, A., Marschik, P. B., Baumeister, H., Dong, F., Hantke, S., Pokorný, F. B., Rathner, E.-M., Bartl-Pokorný, K. D., et al. (2018). The interspeech 2018 computational paralinguistics challenge: Atypical & self-assessed affect, crying & heart beats. In *Interspeech*, pages 122–126.
- Shaver, P., Schwartz, J., Kirson, D., and O’connor, C. (1987). Emotion knowledge: further exploration of a prototype approach. *Journal of personality and social psychology*, 52(6):1061.
- Soleymani, M., Asghari-Esfeden, S., Fu, Y., and Pantic, M. (2015). Analysis of eeg signals and facial expressions for continuous emotion detection. *IEEE Transactions on Affective Computing*, 7(1):17–28.
- Song, S. and Meng, Y. (2015). Detecting concept-level emotion cause in microblogging. In *Proceedings of the 24th International Conference on World Wide Web*, pages 119–120. ACM.
- Stepanov, E. A., Chowdhury, S. A., Bayer, A. O., Ghosh, A., Klasinas, I., Calvo, M., Sanchis, E., and Riccardi, G. (2018). Cross-language transfer of semantic annotation via targeted crowdsourcing: task design and evaluation. *Language Resources and Evaluation*, 52(1):341–364.
- Talmy, L. (2000). *Toward a cognitive semantics*, volume 2. MIT press.
- Tammewar, A., Cervone, A., Messner, E.-M., and Riccardi, G. (2019). Modeling User Context for Valence Prediction from Narratives. In *Proc. Interspeech 2019*, pages 3252–3256.
- Tokuhisa, R., Inui, K., and Matsumoto, Y. (2008). Emotion classification using massive examples extracted from the web. In *Proceedings of the 22nd International Conference on Computational Linguistics-Volume 1*, pages 881–888. Association for Computational Linguistics.
- Xu, R., Hu, J., Lu, Q., Wu, D., and Gui, L. (2017). An ensemble approach for emotion cause detection with event extraction and multi-kernel svms. *Tsinghua Science and Technology*, 22(6):646–659.
- Xu, B., Lin, H., Lin, Y., Diao, Y., Yang, L., and Xu, K. (2019). Extracting emotion causes using learning to rank methods from an information retrieval perspective. *IEEE Access*, 7:15573–15583.
- Yada, S., Ikeda, K., Hoashi, K., and Kageura, K. (2017). A bootstrap method for automatic rule acquisition on emotion cause extraction. In *2017 IEEE International Conference on Data Mining Workshops (ICDMW)*, pages 414–421. IEEE.
- Yu, X., Rong, W., Zhang, Z., Ouyang, Y., and Xiong, Z. (2019). Multiple level hierarchical network-based clause selection for emotion cause extraction. *IEEE Access*, 7:9071–9079.